Data Science in Bioinformatics

week.02.remember.your.statistics.class

Palle Villesen & Thomas Bataillon

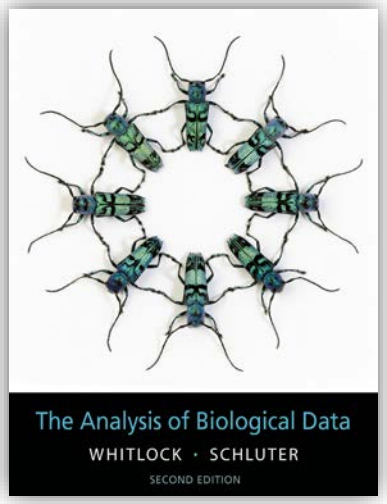

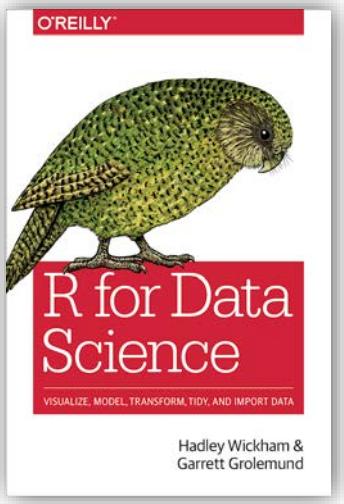

Outline for week 02

- Any questions from last week ?
- Basic statistics
  - Most data are samples
  - Describing data
    - Descriptive statistics of a sample
    - Displaying data
  - Sampling with uncertainty
    - Sampling distributions

Population VS sample

UNKNOWN

- Virtually infinite
- Parameters
- Probability distribution

KNOWN

- n obs: x_1, x_2, …x_n
- Parameter estimates
- Sampling distribution

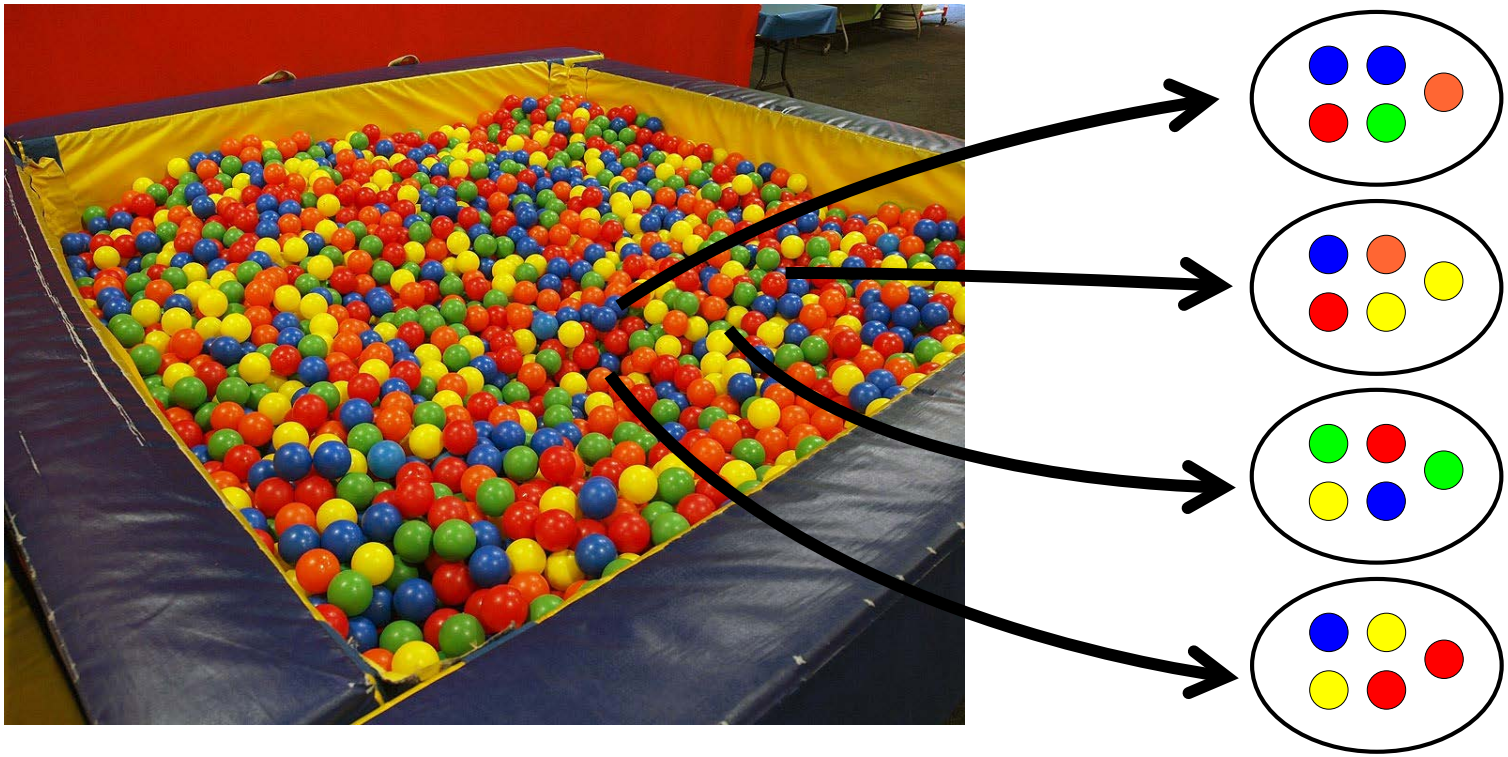

Biased sample: "High-rise syndrome"

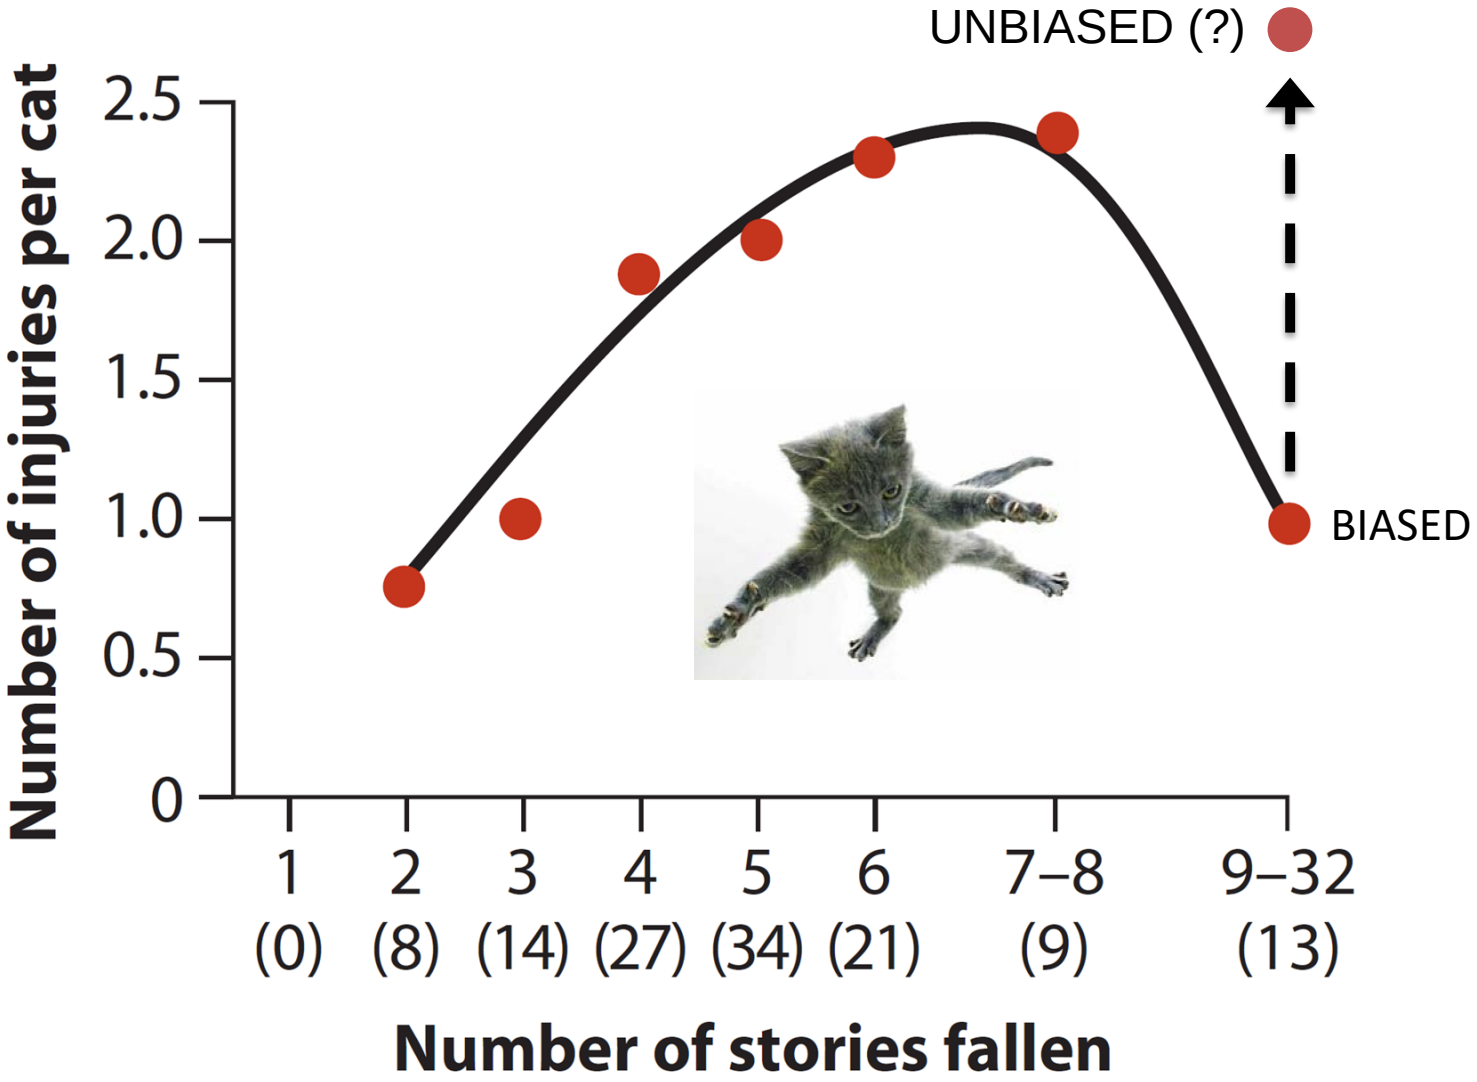

Biased sample

- What if the blue balls are more slippery than the others? ...

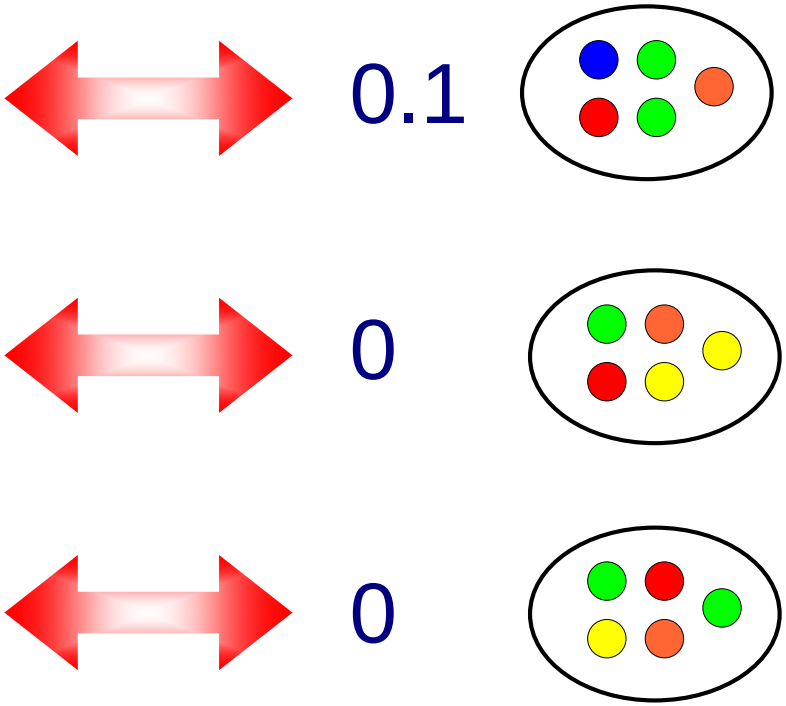

Bullet holes in returning planes in WWII

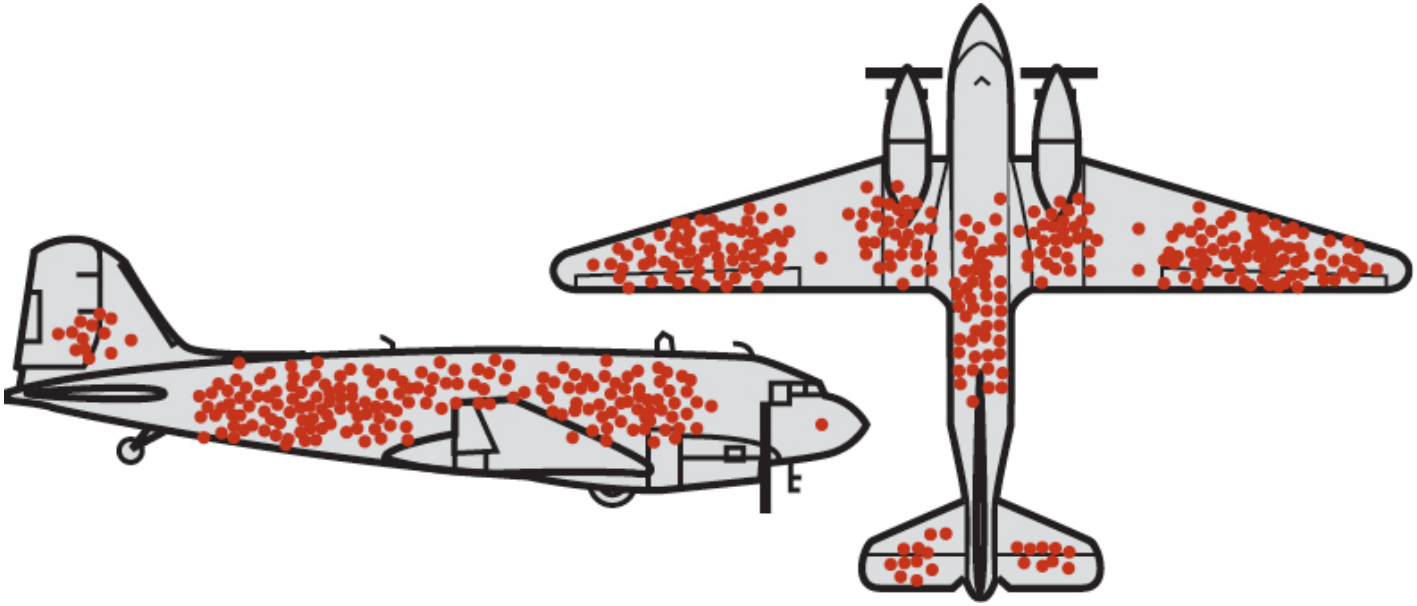

How would you enforce the planes?

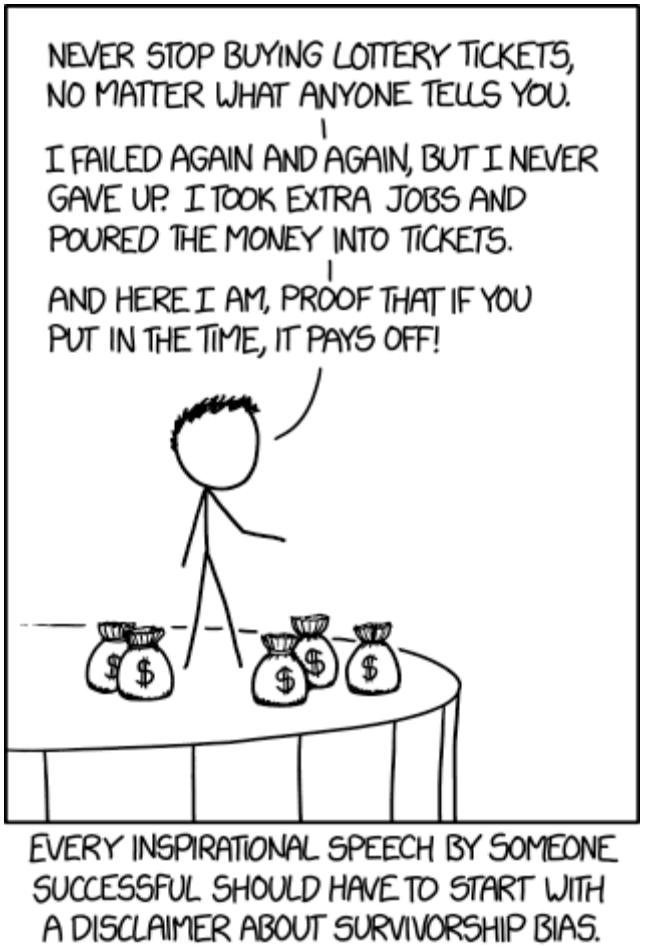

https://xkcd.com/1827/

Displaying data

What is best?

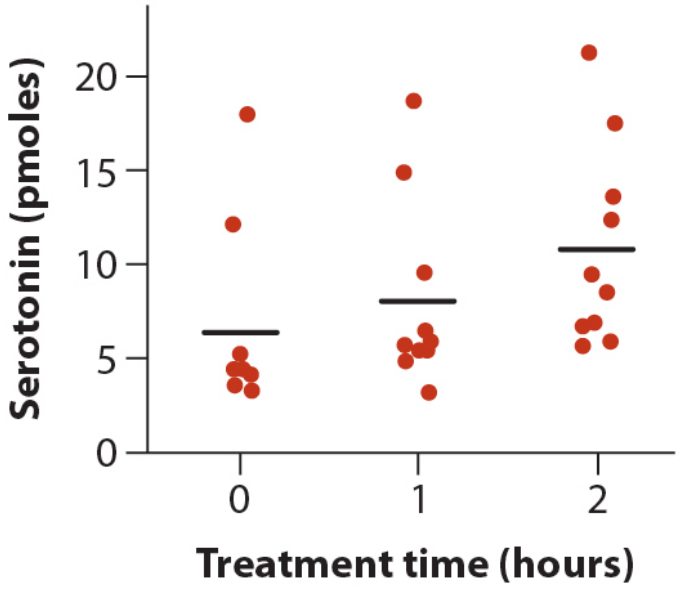
Serotonin (pmoles)
20
15
10
5
0
0
1
2
Treatment time (hours)

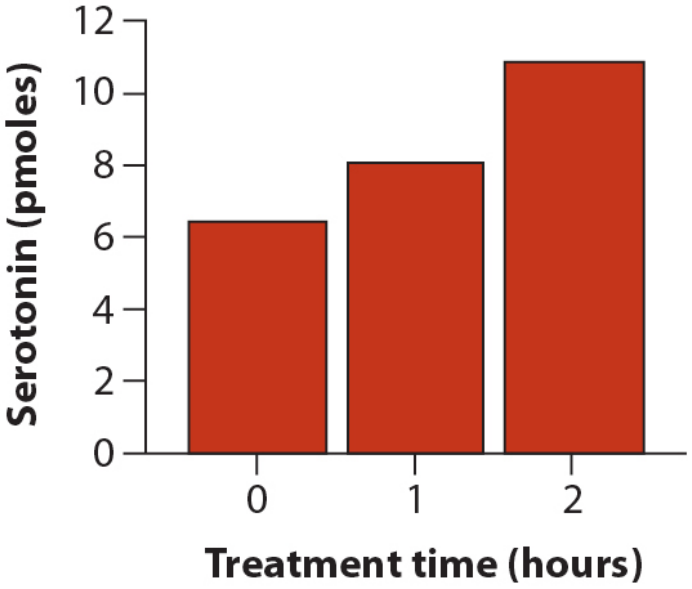
Serotonin (pmoles)
12
10
8
6
4
2
0
0
1
2
Treatment time (hours)

What is best?

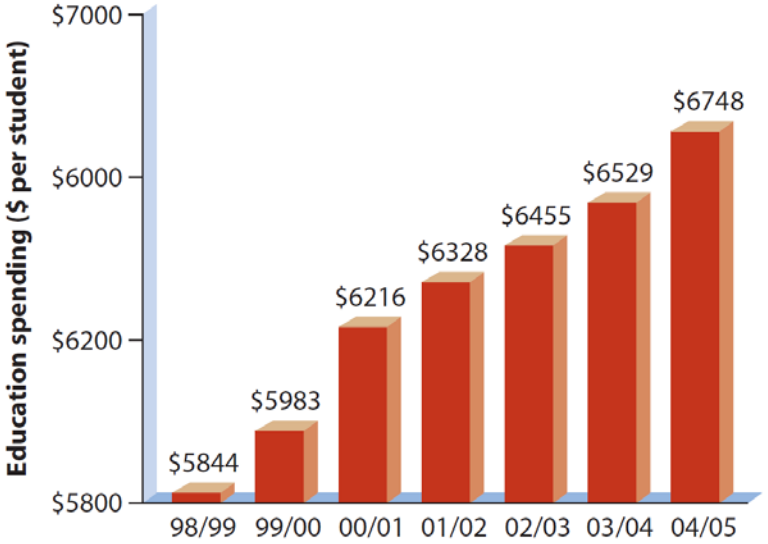
Education spending ($ per student)
$7000
$6000
$6200
$5800
$5844
$5983
$6216
$6328
$6455
$6529
$6748
98/99
99/00
00/01
01/02
02/03
03/04
04/05

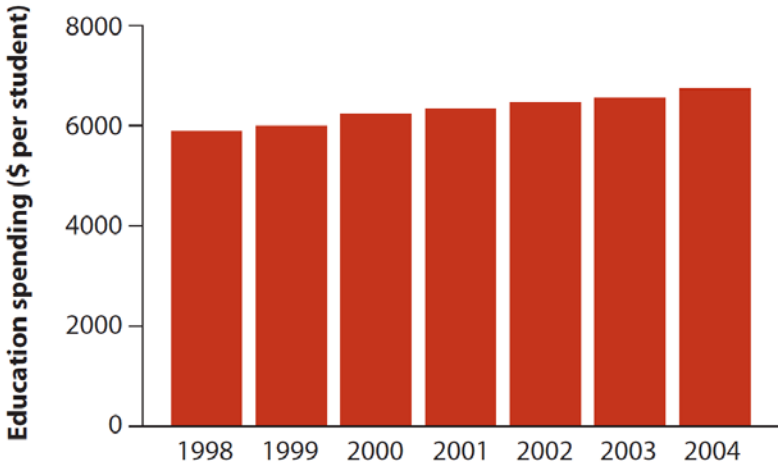
Education spending ($ per student)
8000
6000
4000
2000
0
1998
1999
2000
2001
2002
2003
2004

What is best?

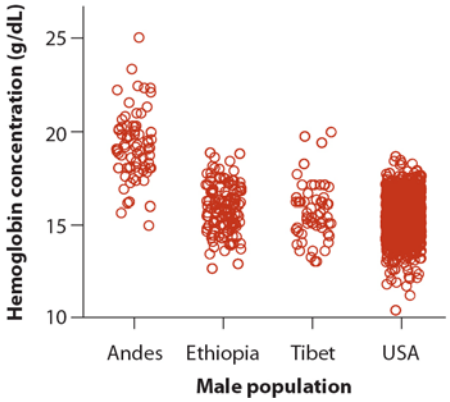
Hemoglobin concentration (g/dL)
25
20
15
10
Andes
Ethiopia
Tibet
USA
Male population

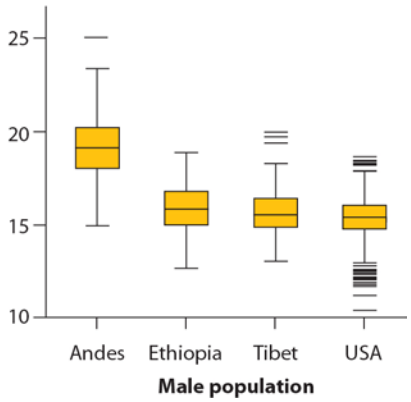
25
20
15
10
Andes
Ethiopia
Tibet
USA
Male population

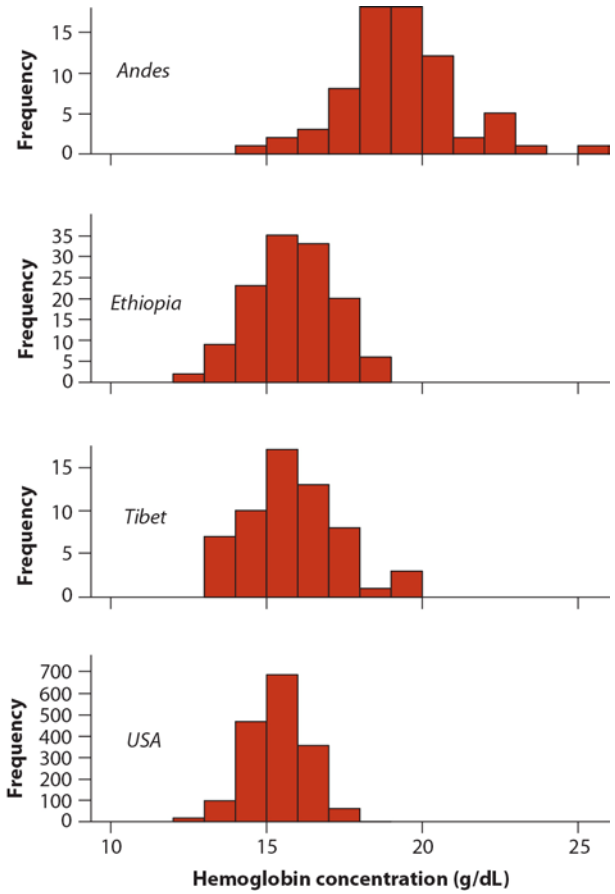
Frequency
Andes
Ethiopia
Tibet
USA
Hemoglobin concentration (g/dL)

Displaying data

- If is really difficult in ggplot, then it is probably a bad idea

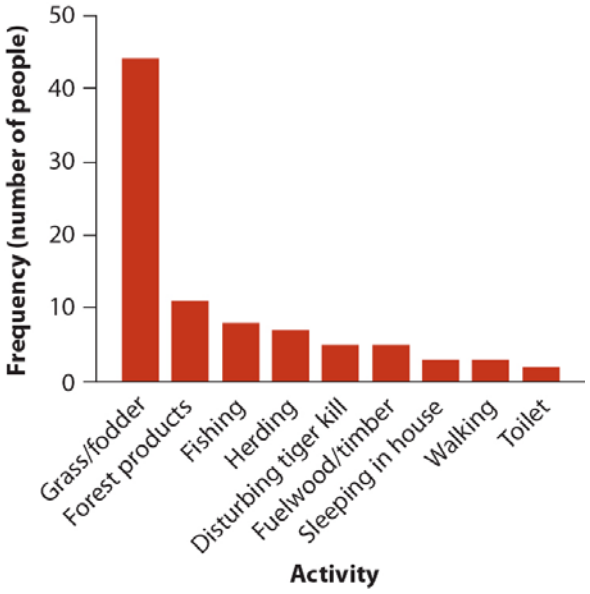

Describing data

- Arithmetic mean = average
- Standard deviation (compares all observations with mean)
- Median = 50% quantile
- IQR (75% quantile – 25% quantile)

- Present a case where one is really bad

Life expectancy is a mean

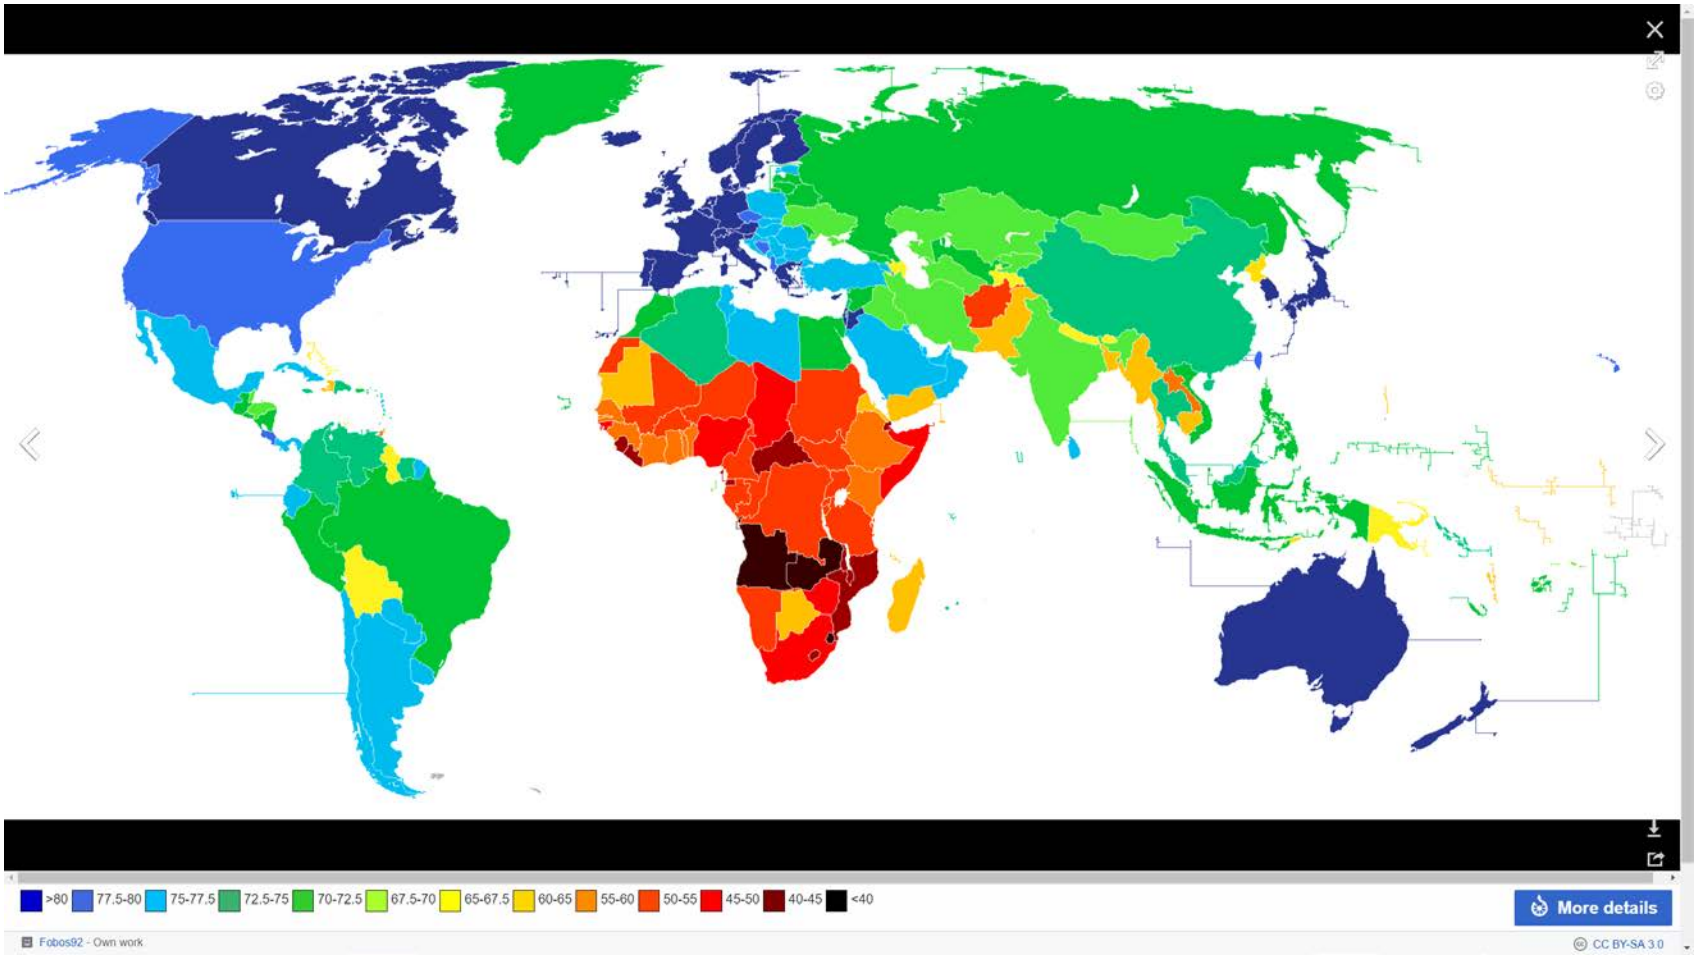

Standard deviation

$$s = \sqrt{\frac{\sum (Y_i - \bar{Y})^2}{n - 1}}$$

Median absolute deviation

$$\mathrm{MAD} = \mathrm{median}(\ |X_i - \mathrm{median}(X)|\ ),$$

Exercise

- Make a small dataset with 3 outliers:

```
x = c(70:90,1000,1100,1200)
```

- Calculate mean and median

```
mean(x)
```

```
median(x)
```

- Calculate sd, iqr and mad

```
sd(x)
```

```
quantile(x)
```

```
mad(x)
```

- Conclusions?

Distribution of the data: boxplots

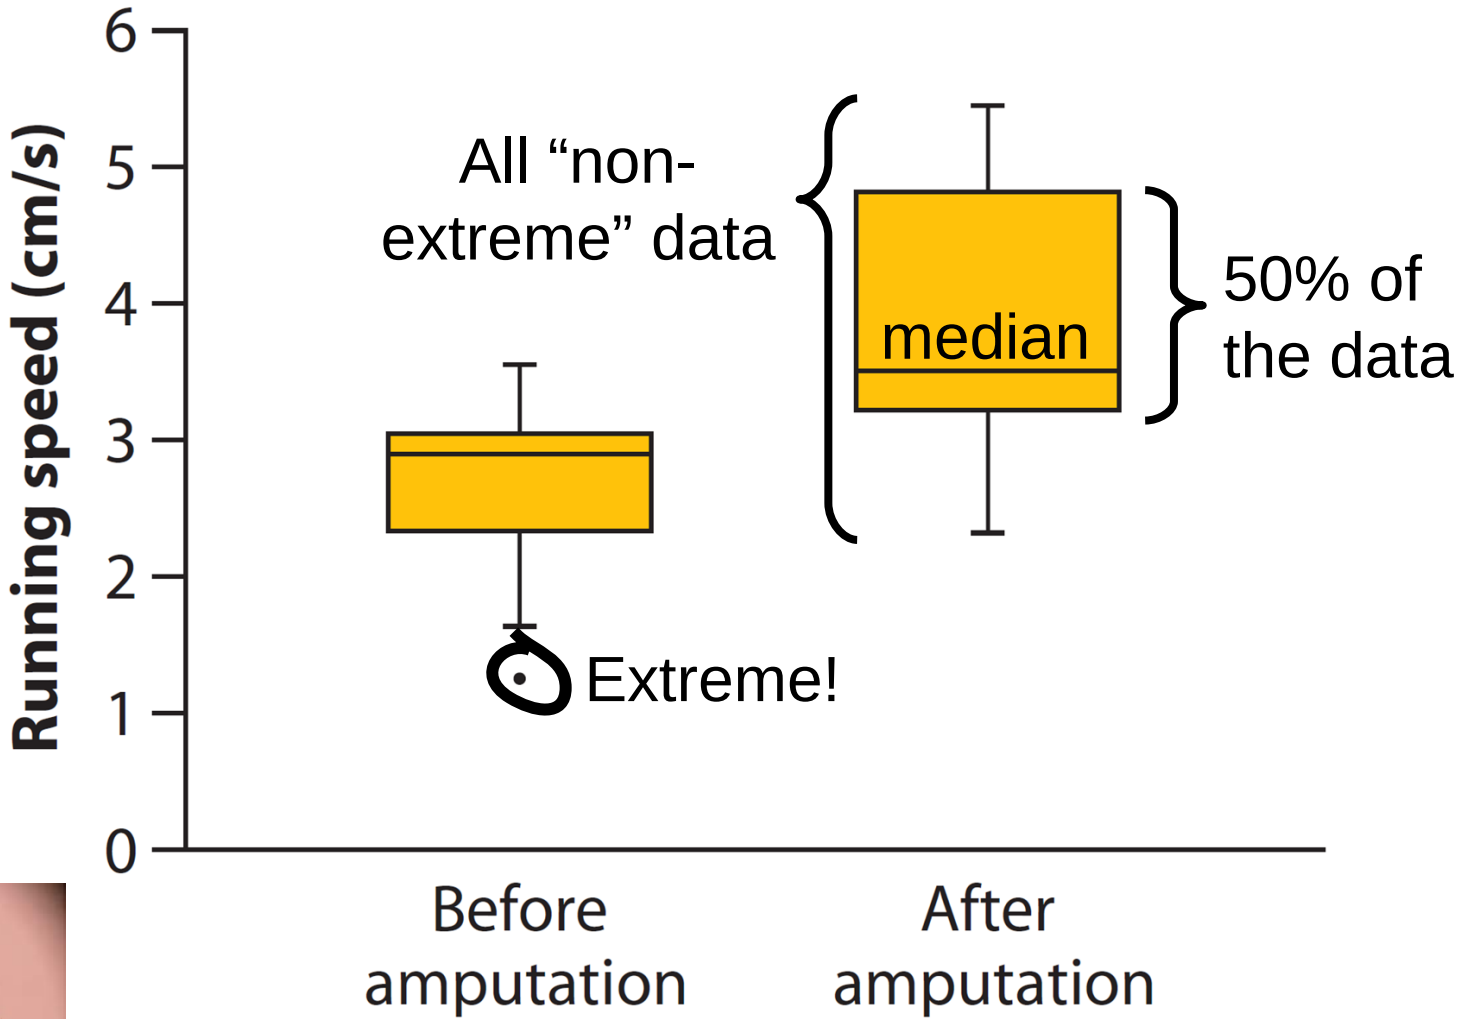

Distribution of the data: histograms

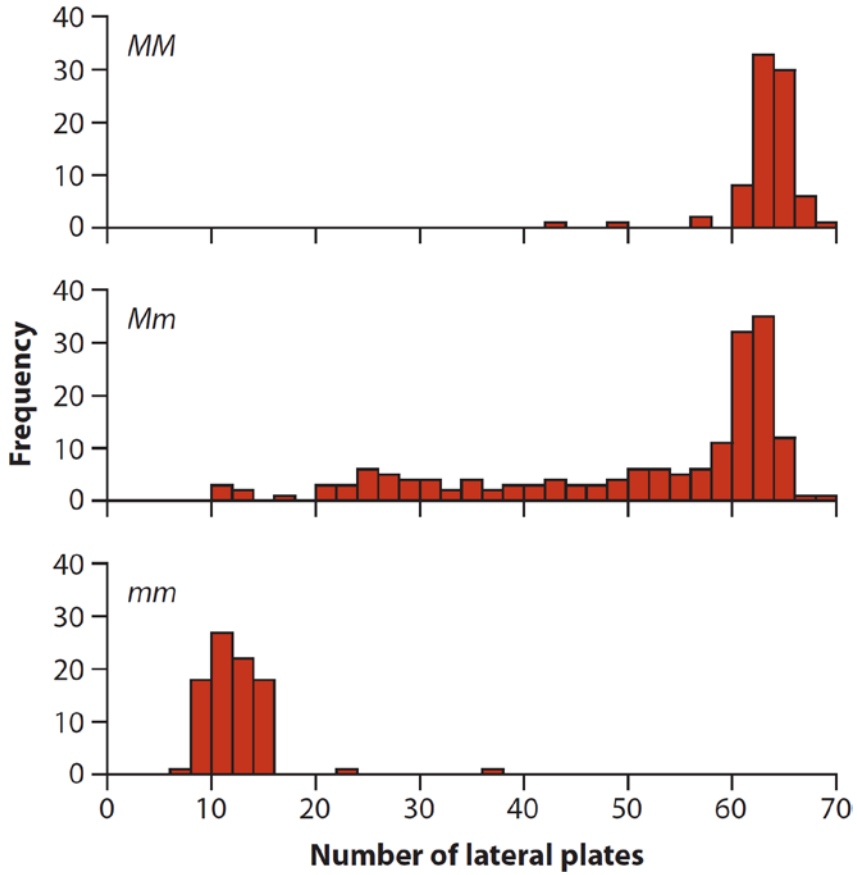

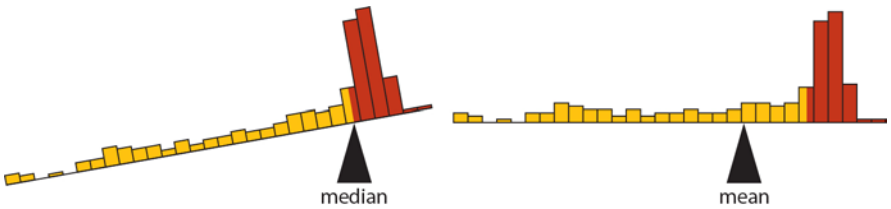

Outliers

- Outliers:
  - Influence mean and sd

- Median and iqr are more robust
  - Median is a single data point
  - iqr only use half of the data

Bodymass of mammals
Discuss and present

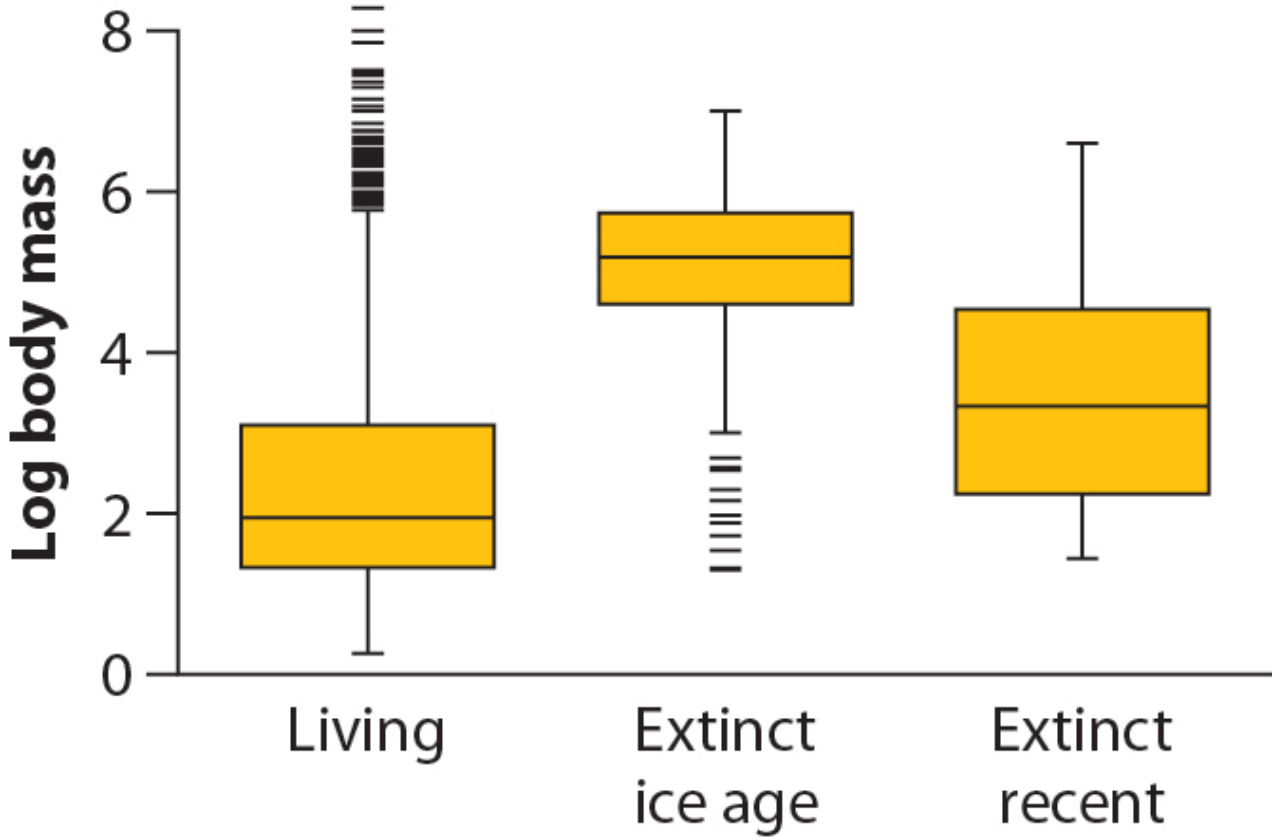

Estimating with uncertainty

Keywords

- We want to say something about the **population**
- But we only have a **sample**
- So the **sample estimate** is different from the true value because of **sampling error**
- **The sampling distribution** is the distribution of estimate from different samples
- **The standard error** is the standard deviation of the sampling distribution
- **Confidence intervals on the estimate**
  - The 2SE rule of thumb
  - Bootstrap (chapter 19)

All genes

```
df <- read_csv(file =
"chap04e1HumanGeneLengths.csv")
```

```
set.seed(0)
dfsub <- df %>% sample_n(size = 100)
```

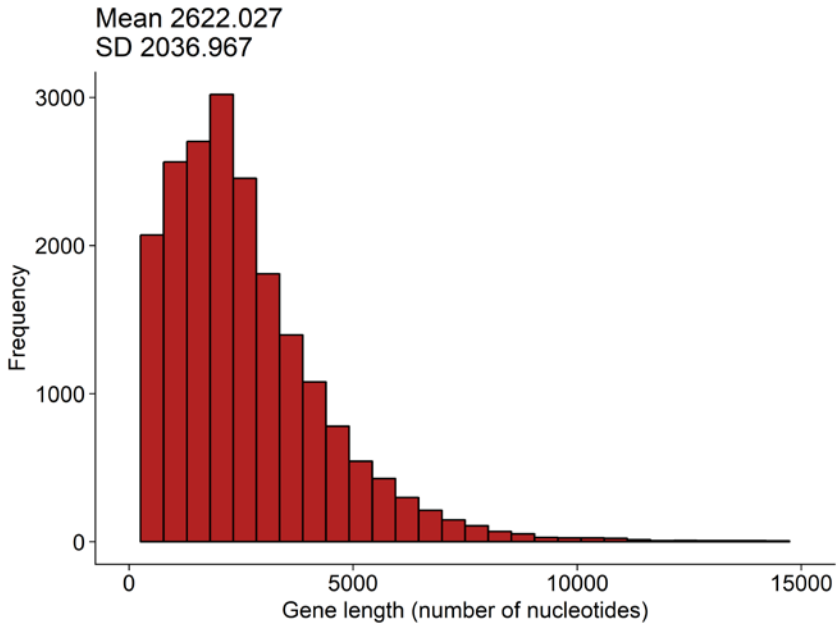

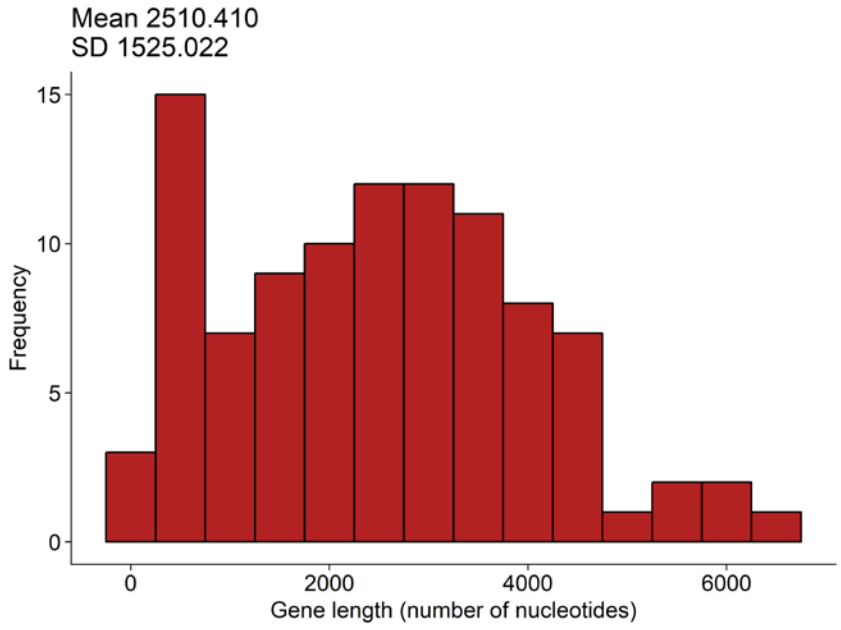

- `df <- read_csv(file="chap04e1HumanGeneLengths.csv")`
- Replicate figure 4.1-3
  - the sampling distribution of the mean for n=100
- Replicate figure 4.1-4
  - the sampling distribution of the mean for n=20, n=100 and n=500
- Calculate standard error from your samples (n=20, n=100, n=500)
- Compare with table 4.2-1

```
set.seed(0)
r <- data.frame() %>% tbl_df()
for (i in 1:10000) {
  dfsub = df %>% sample_n(size = 100)
  r <- rbind(r, data.frame(n = nrow(dfsub), gene.mean = mean(dfsub$geneLength)))
}
```

GROUP WORK